

NIA Python Day3

January 28, 2021

NIA Python Bootcamp UNIT 3 - Wednesday July 19, 2017

1 UNIT 1 review

1. Python ecosystem of tools
2. Jupyter Notebook is code, output and documentation all in one document
3. Type code into cells, and to run them you press Shift-Enter
4. Different data types for different data
5. Tab completion reduces typing, shows you pop-up menu of all the things you can do with that piece of data
6. Operators take one or more input values and turn them into other values *based on the input values type*
7. Converting data from one type to another using the function syntax, e.g., `int()`

2 UNIT 2 Review

1. Exploring data types using the TAB key
2. Python syntax for taking slices of iterables
3. NumPy arrays: basic math operations in 1-D and 2-D (e.g., row-wise and column-wise eman)
4. Subselecting based on a boolean criterion
5. Example: Images as 3-D matrices

3 UNIT 3:

3. PANDAS DataFrames
4. Simple and complex sorting

3.1 PANDAS DataFrame

- `pandas` = [Python Data Analysis Library](#)
- Emulate R's `data.frame` structure.
- Basically a NumPy matrix with
 - Row and column names
 - Can have columns of different types
 - Handles missing data better

3.2 Load the PANDAS package into memory using import()

```
[1]: import pandas as pd
```

3.3 Use PANDAS read_* functions to import data

- There are many functions to import data
- Type `pd.read_` then TAB to see all the import functions

```
[ ]: pd.read_
```

3.4 Read data from file or URL

```
[2]: titanic_data_url = "https://gist.githubusercontent.com/michhar/  
→2dfd2de0d4f8727f873422c5d959fff5/raw/  
→fa71405126017e6a37bea592440b4bee94bf7b9e/titanic.csv"
```

```
[3]: # Optional - you can use R and Python at the same time  
# if you have the Python package rpy2 installed  
%load_ext rpy2.ipynon
```

```
[4]: titanic = pd.read_csv( titanic_data_url )
```

```
[5]: titanic['Cabin'] = titanic['Cabin'].fillna( "" )
```

```
[6]: titanic.head()
```

```
[6]: PassengerId  Survived  Pclass  \  
0             1         0         3  
1             2         1         1  
2             3         1         3  
3             4         1         1  
4             5         0         3
```

```
                                Name      Sex  Age  SibSp  \  
0                Braund, Mr. Owen Harris   male  22.0     1  
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0     1  
2                Heikkinen, Miss. Laina   female  26.0     0  
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)   female  35.0     1  
4                Allen, Mr. William Henry    male  35.0     0
```

```
    Parch      Ticket    Fare Cabin Embarked  
0      0    A/5 21171   7.2500     S  
1      0    PC 17599  71.2833    C85     C  
2      0  STON/O2. 3101282   7.9250     S  
3      0    113803   53.1000   C123     S  
4      0    373450   8.0500     S
```


```
[7]: %>%R -i titanic
library(tidyverse)
```

```
R[write to console]: Attaching packages
                     tidyverse 1.3.0
```

```
R[write to console]: ggplot2 3.3.2
purrr 0.3.4
tibble 3.0.3      dplyr 1.0.1
tidyr 1.1.1      stringr 1.4.0
readr 1.3.1      forcats 0.5.0
```

```
R[write to console]: Conflicts
                     tidyverse_conflicts()
dplyr::filter() masks stats::filter()
dplyr::lag()    masks stats::lag()
```

```
[8]: %>%R
#glimpse( titanic)
ggplot( titanic, aes( Age, Fare ) ) + geom_point( )
```



3.5 Return type is a DataFrame

```
[9]: type(titanic)
```

```
[9]: pandas.core.frame.DataFrame
```

3.6 What did we just load?

```
[10]: titanic.shape
```

```
[10]: (891, 12)
```

3.6.1 Change the number of rows Pandas will display using the `set_option()` function

Use the word `None` if you want to display all of them.

```
[11]: #pd.set_option( 'display.max_rows', 50 )
```

3.6.2 See the first N rows using `.head(N)`

Defaults to first 5

```
[12]: titanic.head(2)
```

```
[12]: PassengerId  Survived  Pclass  \
0              1         0        3
1              2         1        1

                                Name    Sex  Age  SibSp  \
0                        Braund, Mr. Owen Harris   male  22.0    1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0    1

    Parch    Ticket   Fare Cabin Embarked
0      0  A/5 21171   7.2500     S
1      0  PC 17599  71.2833   C85      C
```

3.6.3 See the last N rows using `.tail(N)`

Defaults to last 5.

```
[13]: titanic.tail(1)
```

```
[13]: PassengerId  Survived  Pclass    Name    Sex  Age  SibSp  \
890           891         0        3  Dooley, Mr. Patrick  male  32.0    0

    Parch  Ticket   Fare Cabin Embarked
890      0  370376   7.75     Q
```

3.6.4 See random N rows using `.sample(N)`

```
[14]: titanic.sample(3)
```

```
[14]: PassengerId  Survived  Pclass    Name    Sex  \
825           826         0        3  Flynn, Mr. John   male
702           703         0        3  Barbara, Miss. Saiide female
136           137         1        1  Newsom, Miss. Helen Monypeny female

    Age  SibSp  Parch  Ticket   Fare Cabin Embarked
825  NaN     0      0  368323   6.9500     Q
702  18.0     0      1   2691  14.4542     C
```

```
136  19.0      0      2  11752  26.2833  D47      S
```

3.7 len() return number of observations (rows)

```
[15]: len(titanic)
```

```
[15]: 891
```

3.8 .shape attribute gives the shape

```
[16]: titanic.shape
```

```
[16]: (891, 12)
```

3.9 .describe(): Get basic statistics across all columns

- Detects which columns are quantitative gives descriptive stats for those

```
[17]: titanic.describe()
```

```
[17]:
```

	PassengerId	Survived	Pclass	Age	SibSp \
count	891.000000	891.000000	891.000000	714.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008
std	257.353842	0.486592	0.836071	14.526497	1.102743
min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

3.10 .count() give number of non-empty cells

```
[18]: titanic.count()
```

```
[18]: PassengerId    891
      Survived      891
```

```
Pclass      891
Name        891
Sex         891
Age        714
SibSp      891
Parch      891
Ticket     891
Fare       891
Cabin      891
Embarked   889
dtype: int64
```

3.11 DataFrame row and column headers

- Like a NumPy array, but with column and row headers.
- Enables slicing by headers, and not just indices like with NumPy arrays
- The collection of row headers is stored in the `.index` attribute.
- The collection of column headers is stored in the `.columns` attribute.

```
[19]: titanic.columns
```

```
[19]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
          'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
          dtype='object')
```

```
[20]: titanic.index
```

```
[20]: RangeIndex(start=0, stop=891, step=1)
```

3.12 Get a single column

Two ways to do it:

1. Use the “object-oriented” style of [API](#), i.e., the “dot.”
2. Use the dict style, i.e., key-value style (put the column name into brackets, get the column)
3. The returned data type is a PANDAS Series object, which keeps the index from the DataFrame attached

```
[21]: titanic.columns
```

```
[21]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
          'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
          dtype='object')
```

```
[22]: titanic['Name']
```

```
[22]: 0                                Braund, Mr. Owen Harris
      1      Cumings, Mrs. John Bradley (Florence Briggs Th...
```

```

2                Heikkinen, Miss. Laina
3      Futrelle, Mrs. Jacques Heath (Lily May Peel)
4                Allen, Mr. William Henry
...
886                Montvila, Rev. Juozas
887                Graham, Miss. Margaret Edith
888      Johnston, Miss. Catherine Helen "Carrie"
889                Behr, Mr. Karl Howell
890                Dooley, Mr. Patrick
Name: Name, Length: 891, dtype: object

```

3.13 using .values

```
[24]: titanic['Name'].values[:10]
```

```
[24]: array(['Braund, Mr. Owen Harris',
        'Cumings, Mrs. John Bradley (Florence Briggs Thayer)',
        'Heikkinen, Miss. Laina',
        'Futrelle, Mrs. Jacques Heath (Lily May Peel)',
        'Allen, Mr. William Henry', 'Moran, Mr. James',
        'McCarthy, Mr. Timothy J', 'Palsson, Master. Gosta Leonard',
        'Johnson, Mrs. Oscar W (Elisabeth Vilhelmina Berg)',
        'Nasser, Mrs. Nicholas (Adele Achem)'], dtype=object)

```

3.14 .value_counts()

```
[25]: titanic['Sex']
```

```
[25]: 0      male
      1      female
      2      female
      3      female
      4      male
...
886      male
887      female
888      female
889      male
890      male
Name: Sex, Length: 891, dtype: object

```

```
[26]: titanic['Sex'].value_counts()
```

```
[26]: male      577
      female    314
Name: Sex, dtype: int64

```


3.15 Use .pivot_table() to have a breakdown of the data

3.15.1 For categorical data, use aggfunc='count'

```
[ ]: titanic.pivot_table?
```

```
[27]: titanic.count()
```

```
[27]: PassengerId      891  
      Survived        891  
      Pclass         891  
      Name          891  
      Sex           891  
      Age           714  
      SibSp         891  
      Parch         891  
      Ticket        891  
      Fare          891  
      Cabin         891  
      Embarked      889  
      dtype: int64
```

```
[28]: titanic.pivot_table( values='Survived', index='Pclass',  
                           columns='Sex', aggfunc='count',  
                           margins=True)
```

```
[28]: Sex      female  male  All  
      Pclass  
      1           94   122  216  
      2           76   108  184  
      3          144   347  491  
      All         314   577  891
```

3.15.2 For non-categorical data, can use another statistical measure for aggregation, like mean

```
[29]: titanic.pivot_table( values='Age', index='Sex',  
                           columns='Pclass',  
                           aggfunc='mean', margins=True)
```

```
[29]: Pclass      1      2      3      All  
      Sex  
      female  34.611765  28.722973  21.750000  27.915709  
      male    41.281386  30.740707  26.507589  30.726645  
      All     38.233441  29.877630  25.140620  29.699118
```

3.16 Quick figures

- Execute this Jupyter command `%matplotlib inline` before executing code that makes figures to get Jupyter to render them as output.

```
[ ]: %matplotlib inline
```

3.16.1 Univariate histograms

```
[ ]: titanic['Age'].hist?
```

```
[30]: thing = titanic['Age']
```


```
[31]: type( thing)
```

```
[31]: pandas.core.series.Series
```

```
[ ]: thing.hist?
```

```
[32]: titanic['Age'].hist()
```


```
[32]: <AxesSubplot:>
```



3.16.2 Bivariate scatter plot using the .plot attribute

```
[33]: titanic.plot.scatter( 'Age', 'Fare' )
```

```
[33]: <AxesSubplot:xlabel='Age', ylabel='Fare'>
```



3.17 Missing data in PANDAS

- Represented as np.nan, which stands for “Not A Number”
- NaN has type float
- No missing data representation for an integer!
 - Either convert all to floats to use NaN (recommended!), or
 - Convert values into strings and store empties as "" (less recommended)
 - Establish a “flag” value, e.g., -999 and filter out those before using (not recommended!)

```
[34]: import numpy as np
```

```
[35]: np.nan
```

```
[35]: nan
```

```
[36]: type( np.nan )
```

```
[36]: float
```

3.18 Column data types

- A single column of data within a PANDAS DataFrame is called a Series.
- All values within a Series must be of the same type.
- Use the `.dtypes` attribute to check data types for each column

```
[37]: titanic.head(3)
```

```
[37]:
```

	PassengerId	Survived	Pclass	\		Name	Sex	Age	SibSp	\
0	1	0	3							
1	2	1	1							
2	3	1	3							

		Name	Sex	Age	SibSp	\
0		Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...		female	38.0	1	
2		Heikkinen, Miss. Laina	female	26.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500		S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250		S

```
[38]: titanic.count()
```

```
[38]: PassengerId      891
Survived            891
Pclass              891
Name                891
Sex                 891
Age                 714
SibSp               891
Parch               891
Ticket              891
Fare                891
Cabin               891
Embarked            889
dtype: int64
```

```
[39]: titanic.dtypes
```

```
[39]: PassengerId      int64
Survived            int64
Pclass              int64
Name                object
Sex                 object
Age                 float64
SibSp               int64
```

```
Parch          int64
Ticket         object
Fare          float64
Cabin          object
Embarked       object
dtype: object
```

3.19 Column data types may hint at missing values

When using `pd.read_csv()` and `pd.read_excel()` to load a file from disk, PANDAS will try to pick a data type for a column that makes sense.

- If a float64 (just a fancy float), then missing values in the form of NaN are possible
 - Use `.count()` to count non-empty (non-NaN) values
- If an int64 (just a fancy int), then probably no missing values in that column
- If an object, this almost always means it's a string in there
 - Can represent missing values as "", but `.count()` only works for float data types!

```
[40]: some_emptys = pd.Series( ["","asdf","","","","27",""] )
      print( some_emptys.dtype )
      some_emptys.count()
```

```
object
```

```
[40]: 7
```

3.19.1 Coerce to numeric values using `pd.to_numeric()`

```
[41]: some_emptys = pd.to_numeric( some_emptys, errors='coerce' )
```

```
[42]: some_emptys
```

```
[42]: 0    NaN
      1    NaN
      2    NaN
      3    NaN
      4    NaN
      5    27.0
      6    NaN
      dtype: float64
```

```
[43]: print( some_emptys.dtype )
      some_emptys.count()
```

```
float64
```

```
[43]: 1
```

3.20 Statistics on a DataFrame ignore NaNs (as one might expect)

- In other words, doesn't count missing values as 0

```
[44]: titanic.count()
```

```
[44]: PassengerId      891
      Survived        891
      Pclass          891
      Name            891
      Sex              891
      Age             714
      SibSp           891
      Parch           891
      Ticket          891
      Fare            891
      Cabin           891
      Embarked        889
      dtype: int64
```

```
[45]: titanic['Age'].describe()
```

```
[45]: count      714.000000
      mean       29.699118
      std        14.526497
      min        0.420000
      25%        20.125000
      50%        28.000000
      75%        38.000000
      max        80.000000
      Name: Age, dtype: float64
```

3.21 Using the Seaborn Package for visualization

- Browse [this page](#) to see all the types of nice figures you can make

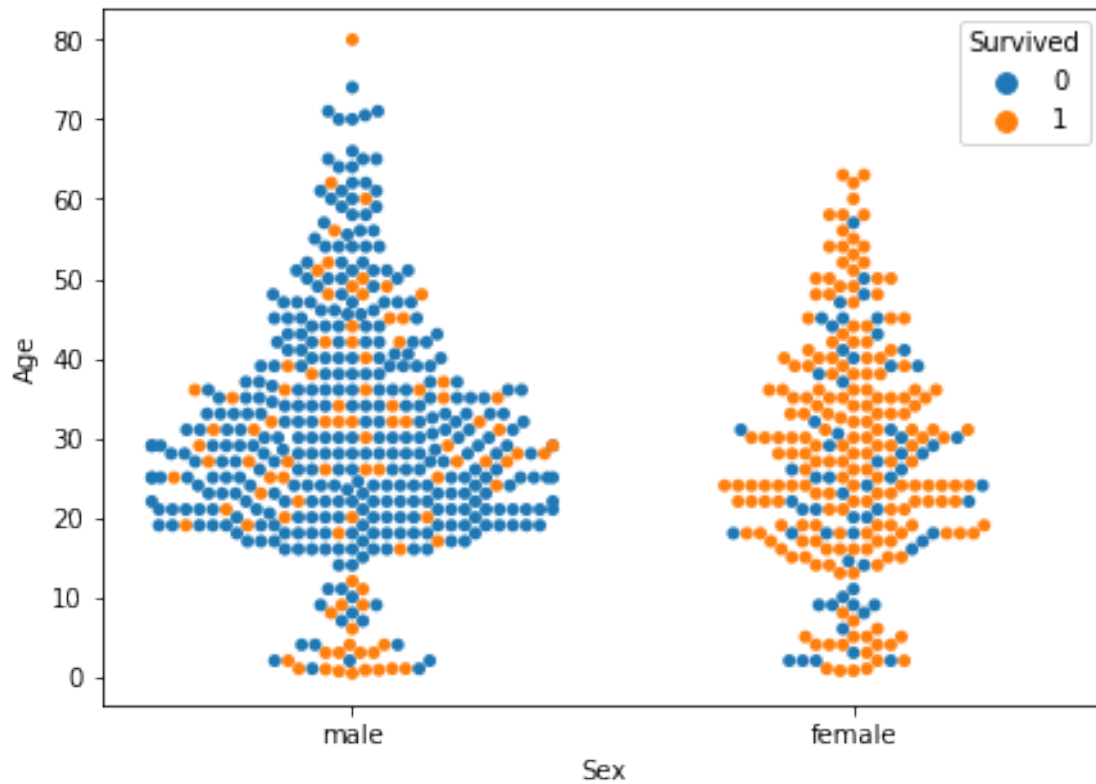
```
[46]: import seaborn as sns
```

```
[47]: import matplotlib.pyplot as plt
```

```
[52]: fig, axis = plt.subplots( figsize=(7,5) )
      sns.swarmplot( x='Sex', y='Age', hue='Survived', data=titanic, ax=axis )

      #fig.savefig( 'testytest.pdf' )
```

```
[52]: <AxesSubplot:xlabel='Sex', ylabel='Age'>
```



```
[53]: type( fig )
```

```
[53]: matplotlib.figure.Figure
```


```
[55]: type( axis)
```

```
[55]: matplotlib.axes._subplots.AxesSubplot
```

```
[ ]: sns.swarmplot?
```

```
[90]: fig, axis = plt.subplots( figsize=(8, 5), dpi=150 )
      sns.swarmplot( x='Pclass', y='Age', hue='Survived', data=titanic, ax=axis)
```

```
[90]: <AxesSubplot:xlabel='Pclass', ylabel='Age'>
```



3.22 Subselecting based on one of the variables

```
[57]: titanic.shape
```

```
[57]: (891, 12)
```

```
[58]: titanic['Sex'].value_counts()
```

```
[58]: male      577
      female    314
      Name: Sex, dtype: int64
```

```
[59]: titanic['Sex'].head()
```

```
[59]: 0    male
      1  female
      2  female
      3  female
      4    male
      Name: Sex, dtype: object
```

```
[60]: titanic['Sex'] == 'male'
```



```
[60]: 0      True
      1      False
      2      False
      3      False
      4      True
      ...
      886     True
      887     False
      888     False
      889     True
      890     True
      Name: Sex, Length: 891, dtype: bool
```

```
[61]: bool_array = titanic['Sex'] == 'male'
```

```
[62]: len(bool_array)
```

```
[62]: 891
```

```
[63]: (titanic['Sex'] == 'male').head()
```

```
[63]: 0      True
      1      False
      2      False
      3      False
      4      True
      Name: Sex, dtype: bool
```

```
[66]: males_only = titanic[ titanic['Sex'] == 'male' ]
```

```
[65]: males_only.head()
```

```
[65]: PassengerId  Survived  Pclass                Name  Sex  Age  \
0             1         0         3    Braund, Mr. Owen Harris  male  22.0
4             5         0         3  Allen, Mr. William Henry  male  35.0
5             6         0         3      Moran, Mr. James  male   NaN
6             7         0         1  McCarthy, Mr. Timothy J  male  54.0
7             8         0         3  Palsson, Master. Gosta Leonard  male   2.0

      SibSp  Parch    Ticket   Fare Cabin Embarked
0         1     0  A/5 21171   7.2500         S
4         0     0   373450   8.0500         S
5         0     0   330877   8.4583         Q
6         0     0    17463  51.8625     E46         S
7         3     1   349909  21.0750         S
```

```
[67]: males_only.shape
```

[67]: (577, 12)

```
[68]: # Boolean selector array have to be the same shape as the array itself!!  
bool_array = [True]*1000
```

```
[74]: titanic[ bool_array ]
```

```
↳ -----  
  
ValueError                                Traceback (most recent call↳  
↳last)  
  
    <ipython-input-74-c15d9c7565a6> in <module>  
    ----> 1 titanic[ bool_array ]  
  
    /usr/local/lib/python3.8/site-packages/pandas/core/frame.py in ↳  
↳__getitem__(self, key)  
    3013         # Do we have a (boolean) 1d indexer?  
    3014         if com.is_bool_indexer(key):  
-> 3015             return self._getitem_bool_array(key)  
    3016  
    3017         # We are left with two options: a single key, and a ↳  
↳collection of keys,  
  
    /usr/local/lib/python3.8/site-packages/pandas/core/frame.py in ↳  
↳_getitem_bool_array(self, key)  
    3060         )  
    3061         elif len(key) != len(self.index):  
-> 3062             raise ValueError(  
    3063                 f"Item wrong length {len(key)} instead of {len(self.↳  
↳index})."  
    3064         )  
  
ValueError: Item wrong length 1000 instead of 891.
```

```
[70]: males_only.shape
```

[70]: (577, 12)

```
[71]: gender_tf = titanic['Sex'] == 'male'
```

```
[72]: gender_tf.shape
```

```
[72]: (891,)
```

```
[73]: males_only.shape
```

```
[73]: (577, 12)
```

```
[77]: females_only = titanic[ titanic['Sex'] == 'female']
```

```
[78]: females_only.shape
```


```
[78]: (314, 12)
```

```
[80]: sns.swarmplot( x='Pclass', y='Age', hue='Survived',  
                  data=males_only)
```

/usr/local/lib/python3.8/site-packages/seaborn/categorical.py:1296: UserWarning:
8.4% of the points cannot be placed; you may want to decrease the size of the
markers or use stripplot.

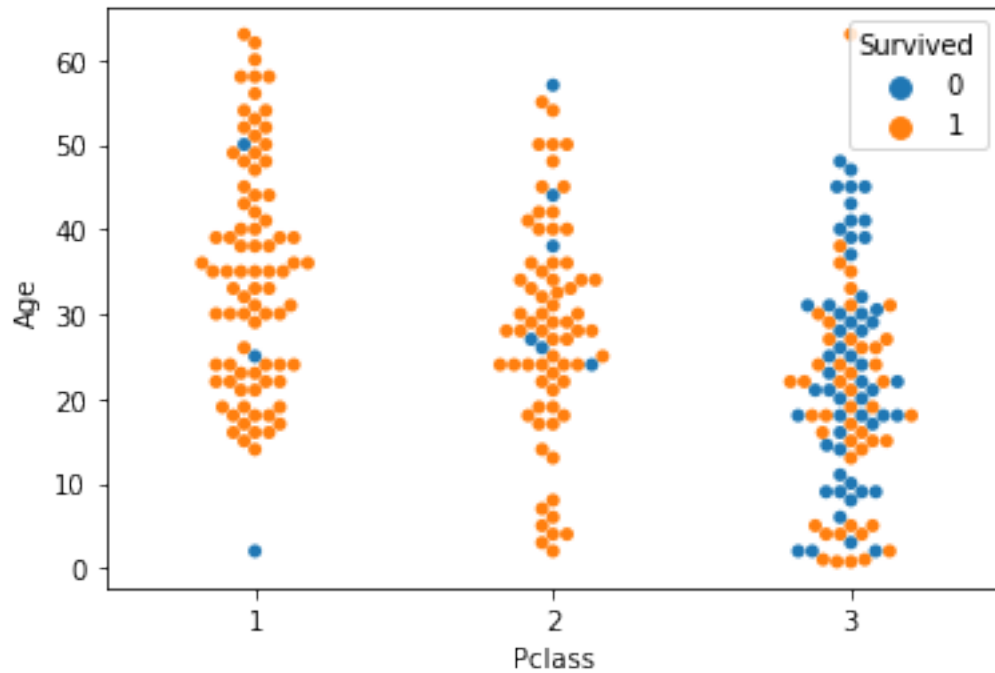
```
warnings.warn(msg, UserWarning)
```

```
[80]: <AxesSubplot:xlabel='Pclass', ylabel='Age'>
```



```
[91]: sns.swarmplot( x='Pclass', y='Age', hue='Survived', data=females_only )
```

```
[91]: <AxesSubplot:xlabel='Pclass', ylabel='Age'>
```



3.23 Slicing by rows and columns using .loc[]

```
[92]: subset = titanic[ titanic['Age'] < 25 ]
```

```
[93]: subset.shape
```

```
[93]: (278, 12)
```

```
[94]: subset = titanic.loc[ titanic['Age'] < 25 ]
```

```
[95]: subset.shape
```

```
[95]: (278, 12)
```

4 Complex sort

```
[96]: age_bool = titanic['Age'] < 10
```

```
[97]: age_bool.value_counts()
```

```
[97]: False    829
      True     62
      Name: Age, dtype: int64
```

```
[98]: class_bool = titanic['Pclass'] == 1
```

```
[99]: class_bool.value_counts()
```

```
[99]: False    675
      True     216
      Name: Pclass, dtype: int64
```

```
[100]: age_class_bool = age_bool & class_bool
```

```
[101]: age_class_bool.value_counts()
```

```
[101]: False    888
      True      3
      dtype: int64
```

```
[103]: titanic.loc[ age_class_bool, 'Age' ]
```

```
[103]: 297    2.00
      305    0.92
      445    4.00
      Name: Age, dtype: float64
```

```
[104]: len(subset)
```

```
[104]: 278
```

4.1 Using .sort_values() for simple or complex sorting

```
[ ]: titanic.sort_values?
```

```
[106]: titanic.shape
```

```
[106]: (891, 12)
```

```
[107]: titanic.sort_values( by=['Pclass', 'Age'] ).head()
```

```
[107]:
```

	PassengerId	Survived	Pclass	Name \	
	305	306	1	1	Allison, Master. Hudson Trevor
	297	298	0	1	Allison, Miss. Helen Loraine
	445	446	1	1	Dodge, Master. Washington
	802	803	1	1	Carter, Master. William Thornton II
	435	436	1	1	Carter, Miss. Lucile Polk

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
305	male	0.92	1	2	113781	151.5500	C22 C26	S
297	female	2.00	1	2	113781	151.5500	C22 C26	S
445	male	4.00	0	2	33638	81.8583	A34	S
802	male	11.00	1	2	113760	120.0000	B96 B98	S
435	female	14.00	1	2	113760	120.0000	B96 B98	S

```
[108]: titanic.sort_values( by=['Pclass','Age'], ascending=False ).head()
```

```
[108]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	\
851	852	0	3	Svensson, Mr. Johan	male	74.0	
116	117	0	3	Connors, Mr. Patrick	male	70.5	
280	281	0	3	Duane, Mr. Frank	male	65.0	
483	484	1	3	Turkula, Mrs. (Hedwig)	female	63.0	
326	327	0	3	Nysveen, Mr. Johan Hansen	male	61.0	

	SibSp	Parch	Ticket	Fare	Cabin	Embarked
851	0	0	347060	7.7750		S
116	0	0	370369	7.7500		Q
280	0	0	336439	7.7500		Q
483	0	0	4134	9.5875		S
326	0	0	345364	6.2375		S